

Why Tensorflow Hub

BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding

Jacob Devlin Ming-Wei Chang Kenton Lee Kristina Toutanova
Google AI Language
{jacobdevlin, mingweichang, kentonl, kristout}@google.com

Abstract

We introduce a new language representation model called **BERT**, which stands for **Bi**directional **E**ncoder **R**epresentations from **T**ransformers. Unlike recent language representation models (Peters et al., 2018a; Radford et al., 2018), BERT is designed to pre-train deep bidirectional representations from unlabeled text by jointly conditioning on both left and right context in all layers. As a result, the pre-trained BERT model can be fine-tuned with just one additional output layer to create state-of-the-art models for a wide range of tasks, such as question answering and language inference, without substantial task-specific architecture modifications.

BERT is conceptually simple and empirically powerful. It obtains new state-of-the-art results on eleven natural language processing tasks, including pushing the GLUE score to 80.5% (7.7% point absolute improvement), MultiNLI accuracy to 86.7% (4.6% absolute improvement), SQuAD v1.1 question answering Test F1 to 93.2 (1.5 point absolute improvement) and SQuAD v2.0 Test F1 to 83.1 (5.1 point absolute improvement).

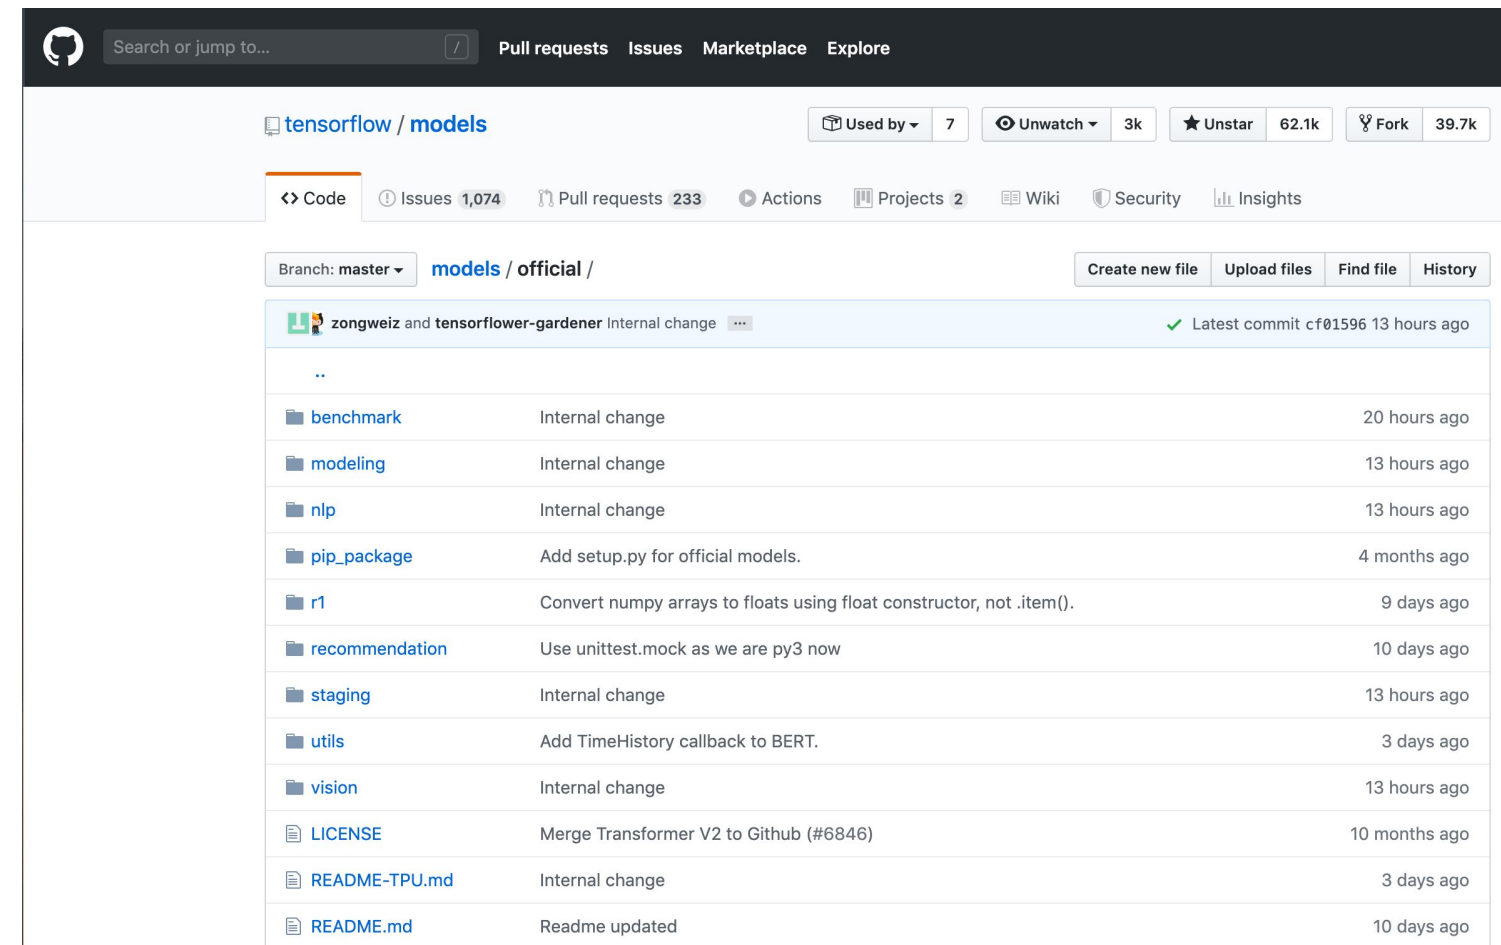

1 Introduction

Language model pre-training has been shown to be effective for improving many natural language processing tasks (Dai and Le, 2015; Peters et al., 2018a; Radford et al., 2018; Howard and Ruder, 2018). These include sentence-level tasks such as natural language inference (Bowman et al., 2015; Williams et al., 2018) and paraphrasing (Dolan and Brockett, 2005), which aim to predict the relationships between sentences by analyzing them holistically, as well as token-level tasks such as named entity recognition and question answering, where models are required to produce fine-grained output at the token level (Tjong Kim Sang and De Meulder, 2003; Rajpurkar et al., 2016).

There are two existing strategies for applying pre-trained language representations to downstream tasks: *feature-based* and *fine-tuning*. The feature-based approach, such as ELMo (Peters et al., 2018a), uses task-specific architectures that include the pre-trained representations as additional features. The fine-tuning approach, such as the Generative Pre-trained Transformer (OpenAI GPT) (Radford et al., 2018), introduces minimal task-specific parameters, and is trained on the downstream tasks by simply fine-tuning *all* pre-trained parameters. The two approaches share the same objective function during pre-training, where they use unidirectional language models to learn general language representations.

We argue that current techniques restrict the power of the pre-trained representations, especially for the fine-tuning approaches. The major limitation is that standard language models are unidirectional, and this limits the choice of architectures that can be used during pre-training. For example, in OpenAI GPT, the authors use a left-to-right architecture, where every token can only attend to previous tokens in the self-attention layers of the Transformer (Vaswani et al., 2017). Such restrictions are sub-optimal for sentence-level tasks, and could be very harmful when applying fine-tuning based approaches to token-level tasks such as question answering, where it is crucial to incorporate context from both directions.

In this paper, we improve the fine-tuning based approaches by proposing BERT: Bidirectional Encoder Representations from Transformers. BERT alleviates the previously mentioned unidirectionality constraint by using a “masked language model” (MLM) pre-training objective, inspired by the Cloze task (Taylor, 1953). The masked language model randomly masks some of the tokens from the input, and the objective is to predict the original vocabulary id of the masked



How do I use it?

Is it safe?

Is it fair?

Is it the latest version?



Google Developers



tfhub.dev

TensorFlow Hub

A comprehensive collection of models



Image



Text



Video



Audio

Before starting


Pre-trained models ready for transfer learning on your own datasets and deployable anywhere you want




TensorFlow
Extended



TensorFlow
Lite



TensorFlow
.JS



Coral

☰ Text embedding

nnlm-en-dim128-with-normalization

Published by: **Google** Updated: 10/24/2019

Token based text embedding trained on English Google News 200B corpus.

NNLM | **Google News**

 Image classification .JS

imagenet/mobilenet_v1_025_128...

Published by: **Google** Updated: 10/24/2019

Imagenet (ILSVRC-2012-CLS) classification with
MobileNet V1 (depth multiplier 0.25).

MobileNet V1 ImageNet (ILSVRC-201...

☰ Text embedding

universal-sentence-encoder-xling/en-fr

Published by: **Google** Updated: 10/24/2019

English and French language-agnostic text encoder.

Transformer

Image pose detection .JS

posenet/mobilenet/float/050

Published by: TensorFlow Updated: 10/24/2019

PoseNet model for pose estimation.

 Image feature vector ..JS

imagenet/mobilenet_v2_100_160...


Published by: **Google** Updated: 10/24/2019

Feature vectors of images with MobileNet V2 (depth multiplier 1.00) trained on ImageNet (ILSVRC-2012-

MobileNet V2 | **ImageNet (ILSVRC-201...**

 Image generator

biggan-deep-512
Published by: DeepMind Updated: 10/24/2019

 **Collection**

universal-sentence-encoder

Published by: **Google** Updated: 10/24/2019

Collection of universal sentence encoders trained on variety of data.

DAN,Transformer


 Image segmentation .JS

bodypix_050

Published by: **TensorFlow** Updated: 10/24/2019

A person segmentation model.

The TensorFlow logo, featuring a stylized orange 'T' inside a white circle, is positioned in the top right corner. Below it, the word 'Publisher' is written in a white sans-serif font. The TensorFlow logo itself is a large, stylized orange 'T' inside a white circle.

 Image feature vector

vae

Published by: **Vtab** Updated: 10/24/2019

Visual representation obtained by training a VAE on ImageNet.

Other **ImageNet (ILSVRC-201...**

☰ Text embedding

nnlm-ja-dim128-with-normalization

Published by: **Google** Updated: 10/24/2019


Token based text embedding trained on Japanese Google News 6B corpus.

NNLM | **Google News**

 Image feature vector

rotation

Published by: **Vtab** Updated: 10/24/2019

 Image pose detection .JS

posenet/mobilenet/quantized/1/100

Published by: **TensorFlow** Updated: 10/24/2019

PoseNet model for pose estimation.

Other

llr-pretrain-adv/latents

Published by: **DeepMind** Updated: 10/24/2019

LLR-pretrained ResNet-152 on ImageNet then adversarially trained for the unrestricted adversarial

Other | Other


 Image classification

[imagenet/pnasnet_large/classification](#)

Published by: **Google** Updated: 10/24/2019

Imagenet (ILSVRC-2012-CLS) classification with
PNASNet-5 (large).

PNASNet-5 (large) | ImageNet (ILSVRC-201...

 Image generator

[compare_gan...](#)

Published by: **Google** Updated: 10/24/2019

ResNet19 trained on CelebA HQ (128x128) (FID: 35.85).

Other | **CelebA HQ**

 Image classification .JS

[imagenet/mobilenet_v2_100_224...](#)

Published by: **Google** Updated: 10/24/2019


Imagenet (ILSVRC-2012-CLS) classification with MobileNet V2 (depth multiplier 1.00).

MobileNet V2 | **ImageNet (ILSVRC-201...**

 Image feature vector

imagenet/mobilenet_v1_100_128...

Published by: **Google** Updated: 10/24/2019

 Image classification

imagenet/mobilenet_v1_025_224...

Published by: **Google** Updated: 10/24/2019

Imagenet (ILSVRC-2012-CLS) classification with
MobileNet V1 (depth multiplier 0.25).

MobileNet V1 | **ImageNet (ILSVRC-201...**

 Image generator

compare_gan/s3gan_20_128x128

Published by: **Google** Updated: 10/24/2019

S3GAN trainend on ImageNet with 20% labels.

Other | **ImageNet (ILSVRC-201...**

 Image feature vector .JS

[imagenet/mobilenet_v1_025_192...](#)

Published by: **Google** Updated: 10/24/2019

Feature vectors of images with MobileNet V1 (depth multiplier 0.25) trained on ImageNet (ILSVRC-2012-

MobileNet V1 **ImageNet (ILSVRC-201...**


☰ Text embedding

bert_en_cased_L-24_H-1024_A-16

Published by: TensorFlow Updated: 10/24/2019

Bidirectional Encoder Representations from Transformers (BERT).

Transformer Wikipedia and BooksC...

 Image classification

unsupervised-adversarial-training...

Published by: **DeepMind** Updated: 10/24/2019

UAT++ adversarially trained WRN-106 (wide residual network) model using 80m@200K unlabeled data and

Other | **CIFAR-10**


 Image classification TFLite
mobilenet_v1_0.25_224
Published by: TensorFlow Updated: 10/24/2019


 Image classification .JS

imagenet/mobilenet_v2_075_224...

Published by: **Google** Updated: 10/24/2019

Imagenet (ILSVRC-2012-CLS) classification with MobileNet V2 (depth multiplier 0.75).

MobileNet V2 ImageNet (ILSVRC-201...

 Image pose detection .JS

posenet/mobilenet/quantized/2/050

Published by: **TensorFlow** Updated: 10/24/2019

PoseNet model for pose estimation.

 Image generator

biggan-deep-256

Published by: **DeepMind** Updated: 10/24/2019

BigGAN-deep image generator trained on 256x256 ImageNet.

Other | ImageNet (ILSVRC-201...

☰ Text embedding

tf2-preview/nnlm-de-dim50-with-...

Published by: **Google** Updated: 10/24/2019

Token based text embedding trained on German Google News 30B corpus.

NNLM | **Google News**

 Image feature vector

imagenet/resnet_v2_152/feature_vector

Published by: **Google** Updated: 10/24/2019

Feature vectors of images with ResNet V2 152 trained on ImageNet (ILSVRC-2012-CLS).

ResNet V2 152 | **ImageNet (ILSVRC-201...**

Image classification TFLite

mobilenet_v1_0.75_192
Published by: TensorFlow Updated: 10/24/2019

[← imagenet/mobilenet_v2_050_96/feature_vector](#)

Problem domain

Image feature vector

Architecture

MobileNet V2

Publisher

Google

Dataset

ImageNet (ILSVRC-2012-CLS)

Format: TF2.0 Saved ModelFine tunable: YesLicense: [Apache-2.0](#) Last updated: 2020-02-20

Model formats

Saved Model

.JS (v1, default)

.JS (v2, default)

.JS (v3, default)

Want to use this model?

To use this model, take a look at the example code, or at [our user guide](#).

You can also try out the associated Colab.

[Copy URL to clipboard](#)[Download Model](#)[Open Colab Notebook](#)

Asset size: 2.62MB

TF2 SavedModel

This is a [SavedModel in TensorFlow 2 format](#). Using it requires TensorFlow 2 (or 1.15) and TensorFlow Hub 0.5.0 or newer.

Overview

MobileNet V2 is a family of neural network architectures for efficient on-device image classification and related tasks, originally published by


- Mark Sandler, Andrew Howard, Menglong Zhu, Andrey Zhmoginov, Liang-Chieh Chen: "[Inverted Residuals and Linear Bottlenecks: Mobile Networks for Classification, Detection and Segmentation](#)", 2018.

Mobilenets come in various sizes controlled by a multiplier for the depth (number of features) in the convolutional layers. They can also be trained for various sizes of input images to control inference speed.

Transfer Learning

Traditional ML


- Isolated, single task learning:
 - Knowledge is not retained or accumulated. Learning is performed w.o. considering past learned knowledge in other tasks



vs

Transfer Learning

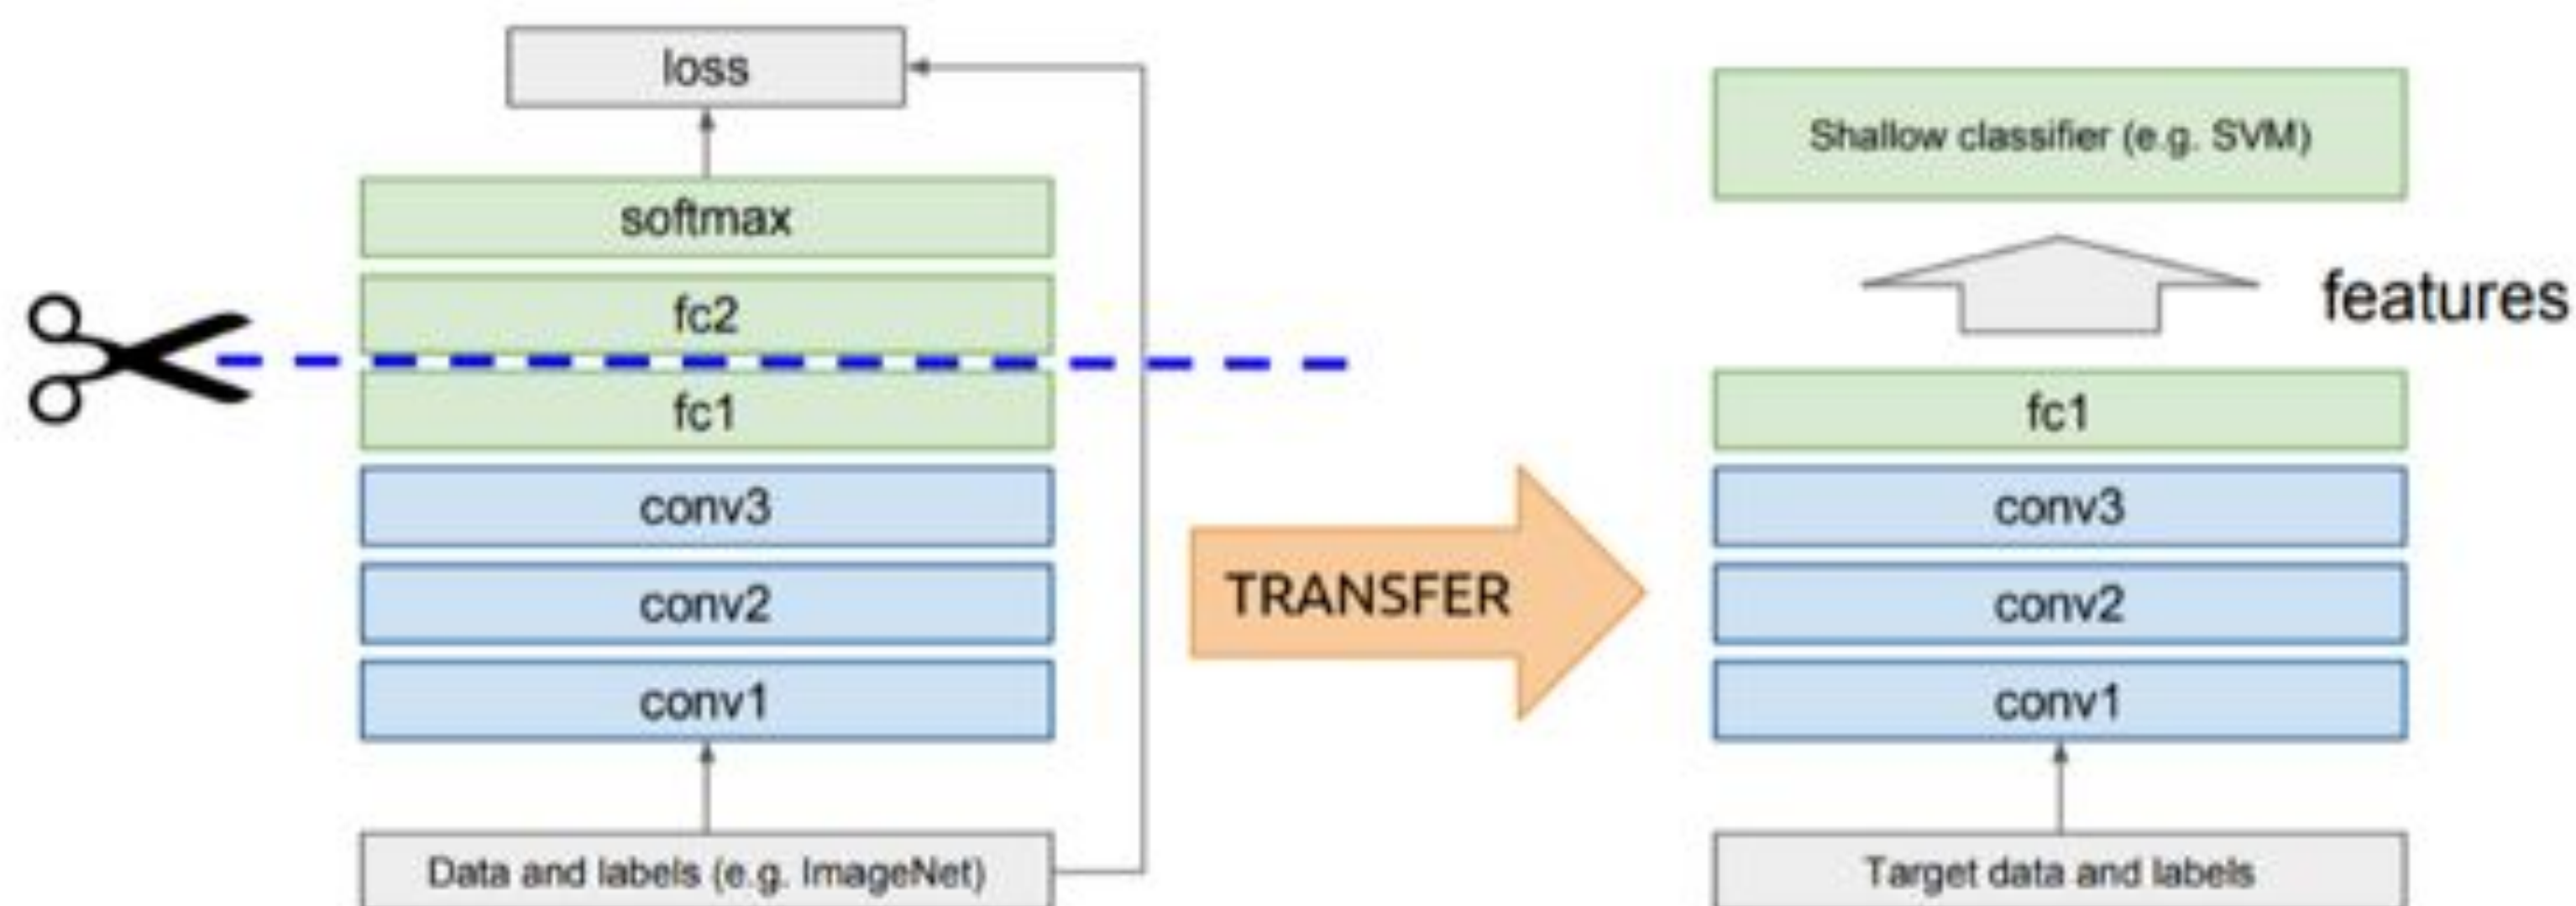
- Learning of a new task relies on the previous learned tasks:
 - Learning process can be faster, more accurate and/or need less training data



Transfer Learning

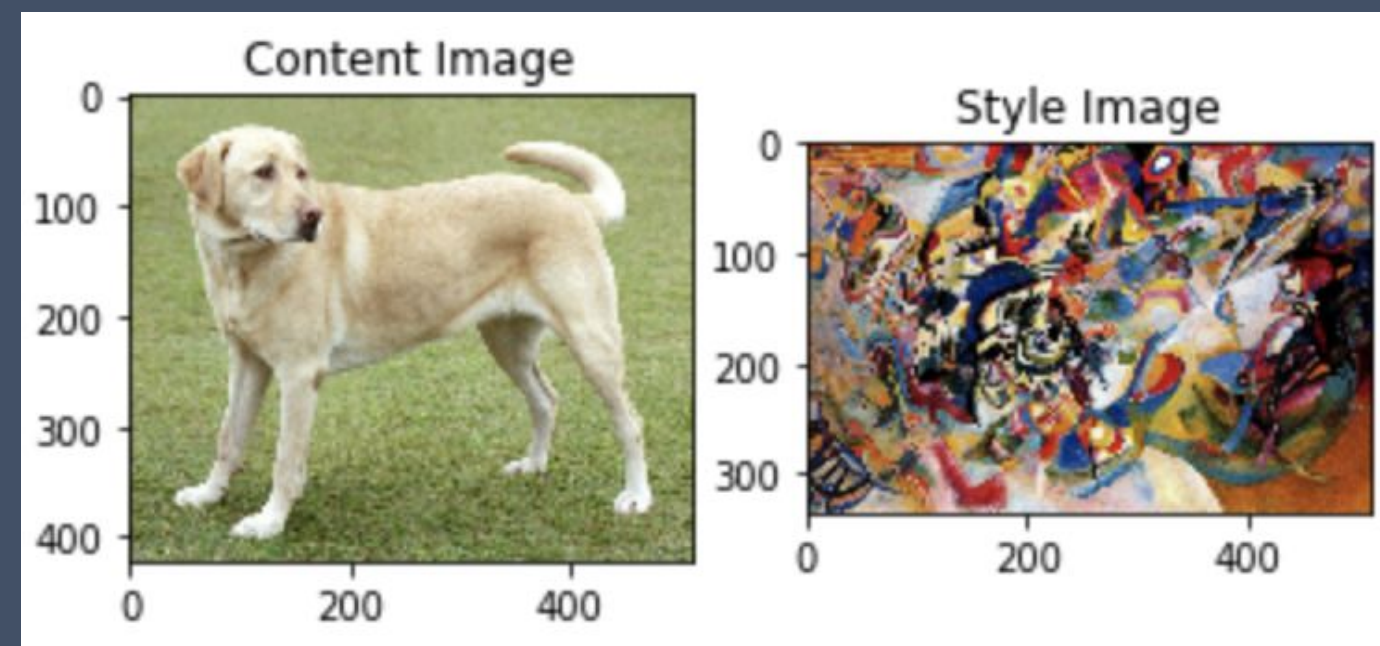
Idea: use outputs of one or more layers of a network trained on a different task as generic feature detectors. Train a new shallow model on these features.

Assumes that $D_S = D_T$



Colab Demo

Style Transfer



```
import tensorflow_hub as hub  
hub_handle = 'https://tfhub.dev/google/magenta/arbitrary-image-stylization-v1-256/1'  
hub_module = hub.load(hub_handle)  
stylized_image = hub_module(tf.constant(content_image), tf.constant(style_image))[0]  
tensor_to_image(stylized_image)
```

tensorflow.org/tutorials/generative/style_transfer

Text Classification

```
import tensorflow as tf
import tensorflow_hub as hub
embedding = "https://tfhub.dev/google/tf2-preview/gnews-swivel-20dim/1"
hub_layer = hub.KerasLayer(embedding, input_shape=[], dtype=tf.string, trainable=True)
```

Text Classification

```
import tensorflow as tf
import tensorflow_hub as hub

embedding = "https://tfhub.dev/google/tf2-preview/gnews-swivel-20dim/1"

hub_layer = hub.KerasLayer(embedding, input_shape=[], dtype=tf.string, trainable=True)

model = tf.keras.Sequential()
model.add(hub_layer)
model.add(tf.keras.layers.Dense(16, activation='relu'))
model.add(tf.keras.layers.Dense(1, activation='sigmoid'))

model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])
history = model.fit(train_data.shuffle(10000).batch(512), epochs=20,
                    validation_data=validation_data.batch(512), verbose=1)
```


What's New

Better Searching for Directories

Improved searching methods for a better experience

TensorFlow Hub

Filters Clear all

Problem domain ▼

Model format

TF.js TFLite Coral

TF Version ?

TF1 TF2

Fine tunable ☐

Architecture ▼

Publisher ▼

Dataset ▼

Language ▼

image classification

Send feedback

Collection

efficientnet

Published by: Google Updated: 02/11/2020

Collection of EfficientNet models for image classification and feature extraction trained on Imagenet (ILSVRC-2012-CLS).

EfficientNet ImageNet (ILSVRC-2012-CLS)

Collection

image

Published by: Google Updated: 02/11/2020

Collection of image models by Google.

ImageNet (ILSVRC-2012-CLS)

Image classification

tf2-preview/inception_v3/classification

Published by: Google Updated: 02/21/2020

[TF2] Imagenet (ILSVRC-2012-CLS) classification with Inception V3.

Inception V3 ImageNet (ILSVRC-2012-CLS)

Image classification

tf2-preview/mobilenet_v2/classification

Published by: Google Updated: 02/21/2020

[TF2] Imagenet (ILSVRC-2012-CLS) classification with MobileNet V2.

MobileNet V2 ImageNet (ILSVRC-2012-CLS)

Image classification

imagenet/mobilenet_v2_100_96/classification

Published by: Google Updated: 02/21/2020

Imagenet (ILSVRC-2012-CLS) classification with MobileNet V2 (depth multiplier 1.00).

MobileNet V2 ImageNet (ILSVRC-2012-CLS)

Image classification

imagenet/mobilenet_v2_100_192/classification

Published by: Google Updated: 02/11/2020

Imagenet (ILSVRC-2012-CLS) classification with MobileNet V2 (depth multiplier 1.00).

MobileNet V2 ImageNet (ILSVRC-2012-CLS)

Expanded Support for TF formats

Tensorflow Lite + Metadata

Additional metadata added to TF-Lite models

Makes deployment on embedded devices easier

TensorFlow.js

New models for face and hand tracking
(published by MediaPipe)

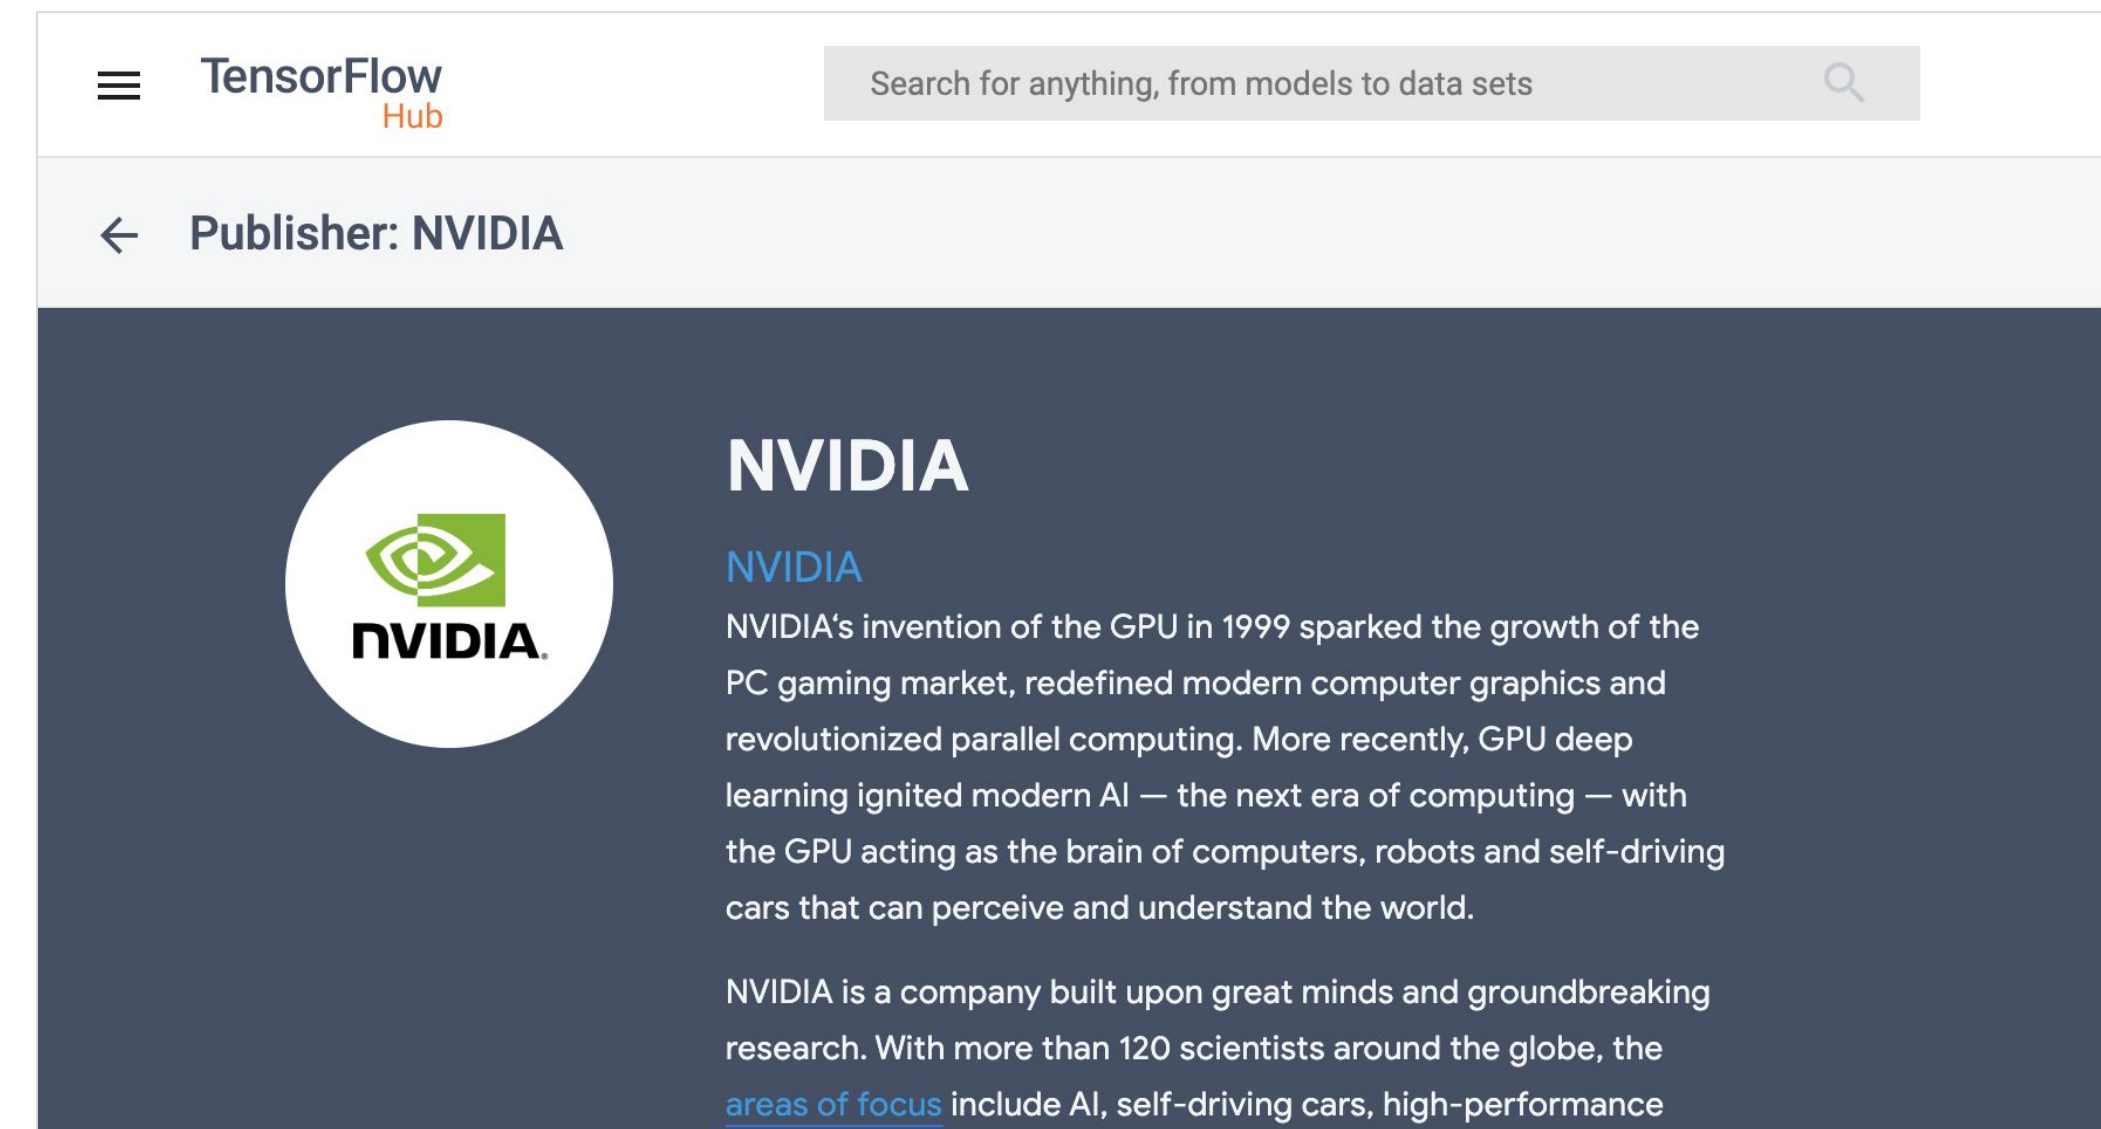
Adding more text models for interactive web apps

With the Help of the Community...

More features will be made compatible with TF after being trained, built and deployed by the community

Some Examples are:

- DeepMind
- Google
- Microsoft AI for Earth
- NVIDIA
- The Metropolitan Museum of Art
- Global Biodiversity Information Facility
- Kaggle
- And more...



Problem Domains

The TensorFlow Hub lets you search and discover hundreds of trained, ready-to-deploy machine learning models in one place. This is made easier by the option to browse by problem domains.

TensorFlow Hub is a repository for machine learning models.

From image classification, text embeddings, audio, and video action recognition, TensorFlow Hub is a space where you can browse trained models and datasets from across the TensorFlow ecosystem. Use it to:

1

Find trained models for transfer learning to save time on training

2

Publish your own models

3

Deploy models on device and in the browser

Search by Problem Domains

On Tensorflow hub you can discover models and collections related to these different problem domains.



Image



Video



Text



Text

Visit <https://tfhub.dev/>

Text

Classification (2) ☐

Embedding (146) ☐

Generation (8) ☐

Language ... (44) ☐

Question ans... (3) ☐

Image

Augmentation (6) ☐

Classificati... (189) ☐

Classification... (1) ☐

Depth estima... (1) ☐

Classifier (5) ☐

Generator (30) ☐

Object dete... (55) ☐

Others (1) ☐

Pose detect... (12) ☐

RNN agent (10) ☐

Video

Classification (6) ☐

Generation (5) ☐

Text (2) ☐

Audio

Embedding (4) ☐

Event classifi... (1) ☐

Pitch extraction (1) ☐

Sample Models

Text Domain - Bidirectional Encoder Representations from Transformers (BERT)

https://www.tensorflow.org/hub/tutorials/bert_experts?hl=en

```
BERT_MODEL = "https://tfhub.dev/google/experts/bert/wiki_books/1" # @param {type: "string"}
["https://tfhub.dev/google/experts/bert/wiki_books/1",
"https://tfhub.dev/google/experts/bert/wiki_books/mnli/1",
"https://tfhub.dev/google/experts/bert/wiki_books/qnli/1",
"https://tfhub.dev/google/experts/bert/wiki_books/qqp/1",
"https://tfhub.dev/google/experts/bert/wiki_books/squad2/1",
"https://tfhub.dev/google/experts/bert/wiki_books/sst2/1",
"https://tfhub.dev/google/experts/bert/pubmed/1", "https://tfhub.dev/google/experts/bert/pubmed/squad2/1"]
```







```
MAX_SEQUENCE_LENGTH = 512
```

```
bert = hub.load(BERT_MODEL)
vocab_path = bert.vocab_file.asset_path.numpy()
tokenizer = tokenization.FullTokenizer(vocab_path, do_lower_case=bert.do_lower_case)
inputs = build_inputs(tokenizer, sentences, MAX_SEQUENCE_LENGTH)
pooled_output, sequence_output = bert(inputs)
```


Sample Models

Image Domain - Object Detection

https://www.tensorflow.org/hub/tutorials/tf2_object_detection?hl=en



Sample Models

Image Domain - Image Classifier (MobileNet V2)

https://www.tensorflow.org/hub/tutorials/tf2_image_retraining?hl=en

```
module_selection = ("mobilenet_v2_100_224", 224)
handle_base, pixels = module_selection
MODULE_HANDLE = "https://tfhub.dev/google/imagenet/{}/feature_vector/4".format(handle_base)
IMAGE_SIZE = (pixels, pixels)
print("Using {} with input size {}".format(MODULE_HANDLE, IMAGE_SIZE))

BATCH_SIZE = 32
```




True label: daisy
Predicted label: daisy

Thanks you

Slide:

<https://bit.ly/2GzrrjL>

